




QUALCOMM SNAPDRAGON X ELITE · HEXAGON V73 HTP

Faraday

Air-Gapped AI Code Assurance & Security Copilot

An offline, air-gapped, on-device AI code assurance and security copilot engineered for the **Qualcomm Snapdragon® X Elite (sc8380xp)** and its **Hexagon™ v73 HTP NPU**. Zero telemetry. Zero egress. Zero compromise.

EXECUTIVE OVERVIEW

The Faraday Philosophy

Named after the **Faraday cage**—an enclosure that blocks electromagnetic fields—**Faraday** forms an impenetrable privacy barrier around your workstation. Source code, internal APIs, and proprietary algorithms never leave your laptop's physical memory.

100% Air-Gapped

Runs with WiFi physically turned off.
No telemetry, no API call-outs, no cloud logins. **0.00 KB egress.**

Hardware-Accelerated

Binds directly to the Qualcomm Hexagon NPU via QNN SDK and ONNX Runtime QNN Execution Provider.

Zero Data Egress

Source code, microservice names, and proprietary algorithms remain in local memory at all times.



Dual-Model Silicon Architecture

Faraday solves the **latency vs. model depth trade-off** by deploying two specialized AI models on the same silicon—each optimized for its role in the code assurance pipeline.

Tier 1 — FaradayCodeAssuranceNeuralNet

AST Neural Risk & Vulnerability Classifier

- **1.41M parameters** — Embedding + Linear Attention + MLP + 3 Classification Heads
- Runtime: ONNX v17 / QNN Context Binary
- File size: **5.66 MB**
- Hexagon NPU latency: **0.21 ms**
- Responsibilities: Sub-ms AST triage, continuous risk score (0–1), 8 CWE classification heads

Tier 2 — Qwen2-7B-Instruct (w4a16)

Generative Code Intelligence Engine

- **7.07B parameters** — Transformer Decoder Block
- Quantization: 4-bit INT4 Weights, FP16 Activations
- Runtime: 4-Part QNN Serialized Binaries
- Memory footprint: **5.05 GB** (fits in LPDDR5x RAM)
- Hexagon NPU speed: **28.4 tokens/second**
- Responsibilities: Root-cause remediation, PEP-257 docstrings, README synthesis

Verifiable Silicon Provider Hierarchy

Faraday maintains strict technical precision—it **never spoofs NPU execution**. The provider hierarchy is transparently reported at runtime, ensuring engineers can always verify where computation is occurring.

① Snapdragon NPU —
QNNExecutionProvider /
libQnnHtp.so

Physical **Hexagon v73 HTP** execution on
Snapdragon X Elite. Sub-millisecond
latency with **0% CPU utilization**. Primary
execution path for both models.

② Neural CPU Fallback —
CPUExecutionProvider

Genuine ONNX tensor forward pass over
code AST token embeddings on x86/ARM
host CPU. Real logits and scores
computed; **CPU status transparently
reported** in all diagnostic output.

③ Deterministic Fallback —
HeuristicRuleBackend

Non-neural fallback utilizing Python AST
visitors and regex patterns when no neural
model artifact is present. Explicitly
declares `is_neural=False` and never
masquerades as NPU.

Qualcomm AI Hub Compilation & Silicon Benchmarks

Compilation Pipeline

FaradayCodeAssuranceNeuralNet: Exported to ONNX v17, compiled to QNN context binary targeting Snapdragon X Elite CRD with w4a16 quantization.

Qwen2-7B-Instruct: Exported as 4-part weight-sharing context binaries totaling **5.05 GB**—fitting entirely in unified LPDDR5x RAM with zero paging.

Silicon Profiling Evidence

Metric	NPU (Tier 1)	NPU (Tier 2)
Physical Core	Hexagon v73 HTP	Hexagon v73 HTP
Quantization	w4a16 / FP16	INT4 / FP16
Latency	0.21 ms	35.2 ms/token
Throughput	470K+ tok/s	28.4 tok/s
CPU Load	0%	0%
Data Egress	0.00 KB	0.00 KB

Complete CLI Command Reference

Command	Purpose	Primary Output
faraday [path]	Full repository audit & documentation generation	REVIEW_REPORT.md
faraday --staged	Scans only Git staged files before commit (<20 ms)	Pre-commit pass/fail
faraday --diff <branch>	Scans modified files against a branch	PR security report
faraday --sarif	Emits OASIS SARIF v2.1.0 for CI/CD	faraday_report.sarif
faraday --ui	Launches interactive local web server	http://localhost:8000
faraday doctor --npu	Audits SoC, Python, QNN DLLs & providers	Silicon health report
faraday --prove	Verifies real tensor execution on silicon	Benchmark certificate
faraday benchmark	Statistical P50/P90/P95/P99 latency benchmarks	benchmark_report.json
faraday --tune	Trains on-device LoRA adapter on local codebase	adapters/*.json
faraday --setup-all	1-Click setup: config + pre-commit + GitHub Actions	.faraday.yml, hooks/

SCENARIOS

Real-World CLI Scenarios



Scenario A–B: Full Audit & Pre-Commit Hook

`faraday ./my_microservice` performs a full AST-chunked audit across Python, JS, and TS. `faraday -staged --fail-on HIGH` intercepts AWS keys or disabled TLS checks in under **20 ms**, aborting commits with exit code 1.



Scenario C–D: PR Diff Scan & CI/CD SARIF

`faraday --diff main` reviews only new PR code, bypassing legacy files. `faraday . --sarif --fail-on HIGH` emits OASIS SARIF v2.1.0 with MITRE CWE descriptors, ingested directly by GitHub Security tab.



Scenario E–F: Visual Dashboard & Silicon Doctor

`faraday --ui` launches an air-gapped SVG dashboard at `localhost:8000` with clickable file explorer and vulnerability highlights. `faraday doctor --npu` and `faraday --prove` verify genuine NPU tensor execution with microsecond latency certificates.



Scenario G–H: Benchmarking & LoRA Fine-Tuning

`faraday benchmark` runs 50+ iterations across sequence lengths (16–128 tokens) for P50/P99 statistics. `faraday . --tune --lora-rank 8` trains LoRA adapters ($W = W_0 + \frac{\alpha}{r}BA$) in under 120 KB files—no cloud exposure.

Static Security Engine & Shannon Entropy

Shannon Entropy Formula

To detect high-entropy secrets—passwords, tokens, encrypted credentials—that bypass keyword filters:

$$H(X) = - \sum_{i=1}^n P(x_i) \log_2 P(x_i) > 4.3 \text{ bits/char}$$

Strings exceeding **4.3 bits/char** are flagged as potential cryptographic material or embedded secrets.

Supported Interceptions (26+ Rules)

- **CWE-89:** SQL Injection via raw string interpolation
- **CWE-79:** DOM XSS via dynamic `innerHTML` assignments
- **CWE-94:** Arbitrary Code Execution via `eval()` / `exec()`
- **CWE-295:** Disabled TLS validation (`verify=False`)
- **CWE-369:** Divide-by-Zero across function blocks
- **CWE-665:** Dangerous mutable default arguments (`def fn(x=[])`)
- **Cloud Keys:** AWS (AKIA...), GitHub (ghp_), OpenAI (sk-...), Anthropic, Google Cloud, HuggingFace, Slack webhooks

36 Automated Tests — 100% Passing in 5.14s

36

Tests Passing

Full verification suite with zero failures

5.14s

Execution Time

Complete suite runtime on local hardware

10

Test Modules

Organized across silicon, security, and CLI layers

0

False Positives

On clean source files — verified by test suite

01

Silicon & NPU Inference (7 tests)

ONNX schema validity, QNN manifest integrity, hardware prover certificates, real tensor execution, heuristic backend labeling, doctor audit, reproducible benchmarks.

02

Model Backend & LoRA Fine-Tuning (8 tests)

Backend selection logic, heuristic rule triage, ONNX export, LoRA forward math
 $h = W_0x + \frac{\alpha}{r}BAx$, parameter efficiency (~15% projection weights), AST corpus dataset, training loop, adapter save/load.

03

Security Scanner & SARIF Governance (6 tests)

Shannon entropy thresholding, AWS/GitHub/OpenAI key interception, clean-file false-positive verification, OASIS SARIF v2.1.0 schema compliance with GitHub URI and CWE rule dictionaries.

04

AST Scanner, Git Utils & Config (10 tests)

Python/JS/TS file discovery with .gitignore, AST function extraction, JavaScript arrow function parsing, staged/diff file detection, default and custom config overrides, --init scaffolding.

05

Report Builder & Vulnerability Regression (5 tests)

REVIEW_REPORT.md synthesis, PEP-257 docstring generation, README architectural synthesis, SQL injection detection, TLS bypass detection, eval/exec detection, XSS innerHTML flagging, hardcoded OpenAI key interception.

Documentation Index & Submission Artifacts

All project documentation, slide presentations, compiled models, and verification artifacts are ready in the repository for review and deployment.



Core Documentation

- **README.md** – Main technical documentation
- **models/README.md** – Silicon model & AI Hub guide
- **.faraday.yml** – Enterprise policy scaffold



Presentation & Dossier

- **Faraday_Pitch_Presentation.html** – 10-slide pitch deck
- **Faraday_Complete_Documentary_Dossier.pdf** – Executive submission dossier



Compiled Model Artifacts

- **faraday_code_assurance.onnx** – ONNX v17 model graph
- **qnn_compile_manifest.json** – Qualcomm compilation manifest
- **models/qwen2-7b-qnn/*.bin** – 4-part QNN serialized binaries (5.05 GB)



Verification Outputs

- **benchmark_report.json** – P50/P90/P95/P99 latency statistics
- **faraday_report.sarif** – OASIS SARIF v2.1.0 CI/CD report
- **36 passing tests** – pytest suite in 5.14s

📄 **Faraday** is production-ready: air-gapped, silicon-verified, and fully documented for enterprise deployment on Qualcomm Snapdragon X Elite platforms.